

University of Cape Town

CSC2002S Assignment PCP1 2023



Dylan Kuming: KMNDYL001

August 13, 2023

Contents

1	Introduction and Methods	3
1.1	Parallelization Approach	3
1.2	Validation	3
1.3	Benchmarking	3
1.4	Machine Architectures	3
1.5	Problems/Difficulties Encountered	3
2	Results	4
2.1	Figures	4
2.2	Discussion	5
2.2.1	Range of grid sizes for optimal performance	5
2.2.2	Maximum speedup and comparison to the ideal	5
2.2.3	Reliability of Measurements	5
2.2.4	Anomalies	5
3	Conclusions	6
A	Appendix	7

List of Figures

1	Speedup versus Grid Size on 2.20 GHz Dual-Core Intel Core i7 machine with Varying Search Densities.	4
2	Speedup versus Grid Size on 1.40 GHz Quad-Core Intel Core i5 machine with Varying Search Densities.	4
3	Table showing data from 2.20 GHz Dual-Core Intel Core i7 machine.	7
4	Table showing data from 1.40 GHz Quad-Core Intel Core i5 machine.	7

Introduction and Methods

1.1 Parallelization Approach

In this assignment, the Java Fork-Join framework is used to speed up the algorithm. The parallelization tactic uses the Monte Carlo algorithm for finding the minimum (the lowest point) of a two-dimensional mathematical function $f(x, y)$ within a specified range.

Given the search density and the grid size, an array of search tasks is initialized. Each search task is designed to find valleys (local minimums) on the terrain. The tasks are divided using the 'ForkJoinPool'. The division occurs recursively: if the number of tasks in a section exceeds the sequential cut-off, it is split into two parts, and each part is processed in parallel. If the number of tasks is below the cut-off, the section is processed sequentially.

An implementation of an optimisation is that if a search reaches a point that a prior search has already visited, it stops as it would trace the same route to the identical local minimum.

1.2 Validation

To validate the parallel Monte Carlo minimization algorithm, first the sequential version will be run on several test terrains, and the results will be stored for comparison. Then, the parallel version will be run on the same terrains, and if the outputs match, this ensures that the parallel program is correct. The Rosenbrock function will be used for this process as it has a known minimum of height 0 at points (1,1). Multiple runs will be executed to identify and remove outliers for more accurate results.

1.3 Benchmarking

Benchmarking will be performed using a timer that records the time between the start and end of the search tasks. The start time is recorded right before initiating the parallel search and the end time is recorded right after completing the search. The performance of the algorithm will be assessed using various grid sizes and search densities in order to ensure sufficient coverage. Tests will be run on two different machine architectures, and the benchmark will be repeated multiple times to ensure consistency.

1.4 Machine Architectures

The algorithms will be tested on two different machine architectures:

1. Machine A: This architecture features a 2.20 GHz Dual-Core Intel Core i7 processor.
2. Machine B: This architecture features a 1.40 GHz Quad-Core Intel Core i5 processor.

1.5 Problems/Difficulties Encountered

Deciding the appropriate sequential cut-off for dividing tasks is very important. If the sequential cut-off is set too low, then there will be excessive task division leading to high overhead from thread management. While setting the sequential cut-off too high will result in under-utilization of available cores as tasks are processed sequentially even when parallelism would be beneficial. Through trial and error, a sequential cutoff of 10000 was chosen.

Additionally, since the parallel solutions share a TerrainArea object, a potential race condition arises when threads write to the same grid spot. This will however be ignored since the race condition is minor, and protecting against it might result in the program being slower, without much benefit gained.

Results

2.1 Figures

Two graphs are shown below in which figure 1 represents the speedup of the parallel program on a 2.20 GHz Dual-Core Intel Core i7 machine, while figure 2 represents the speedup of the parallel program on a 1.40 GHz Quad-Core Intel Core i5 machine.

The speedup versus grid size is plotted for different search densities of 0.1, 0.25, and 0.5. The terrain is kept at a constant size with parameters of $x_{\min} = -10$; $x_{\max} = 10$; $y_{\min} = -10$; $y_{\max} = 10$ and a `SEQUENTIAL_CUTOFF` = 1000. The grid size ranges from 500×500 to 10000×10000 .

The speedup is calculated using the formula:

$$\text{Speedup} = \frac{T_{\text{Serial}}}{T_{\text{Parallel}}}$$

and is plotted using a logarithmic trend-line to ignore outliers which affect the results. Additionally, the tests for T_{Serial} and T_{Parallel} were run 6 different times, and the quickest time was recorded and used to calculate the speedup.

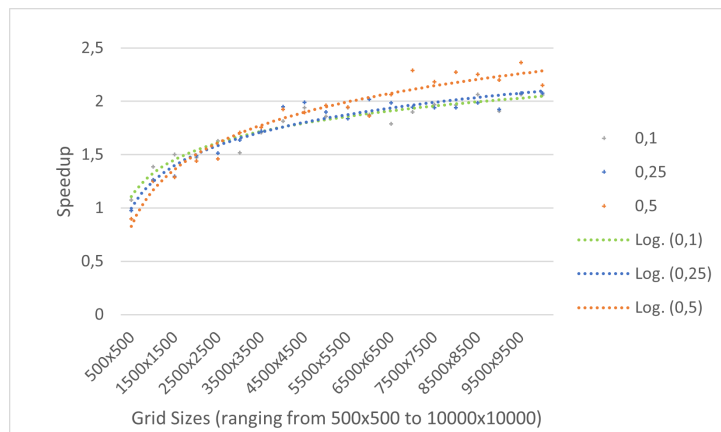


Figure 1: *Speedup versus Grid Size on 2.20 GHz Dual-Core Intel Core i7 machine with Varying Search Densities.*




Figure 2: *Speedup versus Grid Size on 1.40 GHz Quad-Core Intel Core i5 machine with Varying Search Densities.*

2.2 Discussion

2.2.1 Range of grid sizes for optimal performance

The parallel program performs exceptionally well for medium to large grid sizes. For small grid sizes, the overhead of parallelization seems to offset the benefits. This is evident in the speedup graphs, where the speedup for small grid sizes is closer to 1, indicating near-sequential performance. As the grid size increases, the program starts to show significant speedups, reaching its peak at medium grid sizes. Beyond a certain point, however, there seems to be a saturation in speedup due to the inherent limits of parallelization and possibly memory bottlenecks.

For the 2.20 GHz Dual-Core Intel Core i7, the parallel program performs notably well for grid sizes from 4000x4000 and above, reaching peak performance around 7000x7000, and thereafter remaining somewhat constant until 10000x10000.

For the 1.4 GHz Quad-Core Intel Core i5, optimal performance is seen for grid sizes from around 6500x6500 and remaining relatively constant until 10000x10000.

2.2.2 Maximum speedup and comparison to the ideal

On the 2.20 GHz Dual-Core Intel Core i7, the maximum speedup achieved was approximately 2.362 when the search density was 0.5. This is greater than the ideal expected speedup of 2, considering it's a dual-core processor. This suggests super-linear speedup, possibly due the fact that while distributing tasks between the two cores, the data might be better located in cache memory, ensuring faster access and reducing wait times.

On the 1.4 GHz Quad-Core Intel Core i5, the highest speedup observed was approximately 1.768, which is a bit below the ideal expected speedup of 4 for a quad-core machine. The sub-optimal performance could be because the algorithm's parallel version might not be fully optimized to leverage the quad-core system's potential.

2.2.3 Reliability of Measurements

The measurements are relatively reliable. Each experiment was repeated 6 times to minimize the effects of outliers and the speedup is plotted using a logarithmic trend-line to further ignore outliers which effect the results. Furthermore, external factors like other running processes were controlled to the best extent possible to ensure consistent results.

2.2.4 Anomalies

There are instances, especially on the 2.20 GHz Dual-Core Intel Core i7 with search density 0.25, where the speedup drops slightly as the grid size increases, such as between the 4000x4000 and 5000x5000 grid sizes. This is perhaps due to the overheads associated with parallelization, which might occasionally offset the advantages of parallel execution.

On the 1.4 GHz Quad-Core Intel Core i5, for search density 0.25, there's a dip in speedup for the 5500x5500 grid size. This anomaly might be due to system-specific interruptions or inefficiencies in workload distribution among the four cores.

Conclusions

Using parallelization in Java for this problem demonstrated significant potential:

1. **Performance Gains:** Both the 2.20 GHz Dual-Core Intel Core i7 and the 1.4 GHz Quad-Core Intel Core i5 showed improved speeds over the serial version with the 2.20 GHz Dual-Core Intel Core i7 machine even surpassing the expected dual-core speedup.
2. **Efficiency and Overheads:** While the parallel approach improved the performance of the algorithm, it also introduced complexities, such as race conditions. These overheads were ignored for our problem but could be more problematic in other algorithms and should always be accounted for.
3. **Development Complexity:** Parallelization increases code complexity, making development and debugging more challenging. This added complexity is however justified by the performance improvements.

In summary, for this problem, parallelization in Java is beneficial and offers performance advantages. However, the performance gains must be weighed up against the added complexities when considering this approach.

Appendix

Grid size	sequential cutoff = 0.1			sequential cutoff = 0.25			sequential cutoff = 0.5		
	serial (ms)	parallel (ms)	speedup	serial (ms)	parallel (ms)	speedup	serial(ms)	parallel(ms)	speedup
500x500	44	41	1,0731707	43	44	0,977273	52	58	0,896552
1000x1000	133	96	1,3854167	178	141	1,262411	169	135	1,251852
1500x1500	360	240	1,5	275	212	1,29717	303	235	1,289362
2000x2000	484	329	1,4711246	516	347	1,487032	557	387	1,439276
2500x2500	818	503	1,6262425	794	525	1,512381	840	575	1,46087
3000x3000	1107	729	1,5185185	1184	723	1,637621	1334	782	1,705882
3500x3500	1615	946	1,7071882	1716	998	1,719439	1834	1045	1,755024
4000x4000	2206	1218	1,8111658	2461	1264	1,946994	2792	1452	1,922865
4500x4500	3091	1593	1,9403641	3217	1618	1,988257	3500	1847	1,894965
5000x5000	3765	2025	1,8592593	4002	2109	1,897582	4384	2236	1,960644
5500x5500	4676	2412	1,9386401	4986	2712	1,838496	5477	2819	1,942888
6000x6000	5617	2997	1,8742075	5886	2917	2,017827	6557	3524	1,86067
6500x6500	6551	3664	1,7879367	6946	3501	1,984005	8483	4109	2,064493
7000x7000	7796	4104	1,8996101	8149	4191	1,944405	10642	4645	2,291066
7500x7500	8986	4567	1,9675936	9427	4855	1,94171	12014	5505	2,18238
8000x8000	10428	5304	1,9660633	11006	5671	1,940751	14438	6355	2,271912
8500x8500	12368	5991	2,06443	12728	6416	1,983791	16242	7213	2,251768
9000x9000	13554	7107	1,9071338	14294	7433	1,923046	17871	8122	2,20032
9500x9500	15973	7748	2,0615643	17483	8415	2,0776	20893	8843	2,36266
10000x10000	17847	8662	2,0603787	18178	8750	2,077486	21243	9878	2,150537

Figure 3: Table showing data from 2.20 GHz Dual-Core Intel Core i7 machine.

Grid size	sequential cutoff = 0.1			sequential cutoff = 0.25			sequential cutoff = 0.5		
	serial (ms)	parallel (ms)	speedup	serial (ms)	parallel (ms)	speedup	serial(ms)	parallel(ms)	speedup
500x500	40	51	0,784314	37	42	0,880952	40	55	0,727273
1000x1000	124	143	0,867133	125	147	0,85034	136	138	0,985507
1500x1500	376	268	1,402985	373	315	1,184127	333	324	1,027778
2000x2000	563	389	1,447301	601	426	1,410798	601	440	1,365909
2500x2500	855	570	1,5	1001	621	1,611916	961	647	1,485317
3000x3000	1189	894	1,329978	1218	847	1,438017	1352	885	1,527684
3500x3500	1691	1200	1,409167	1775	1226	1,447798	1831	1222	1,498363
4000x4000	2131	1322	1,611952	2288	1588	1,440806	2575	1587	1,622558
4500x4500	2863	1937	1,478059	3042	1927	1,57862	3375	1960	1,721939
5000x5000	3411	2257	1,511298	3931	2325	1,690753	4445	2702	1,645078
5500x5500	4234	2791	1,517019	4453	2817	1,58076	4832	2932	1,648022
6000x6000	5239	3276	1,599206	5789	3490	1,658739	5715	3482	1,641298
6500x6500	6172	3989	1,547255	6877	3889	1,768321	7255	4294	1,689567
7000x7000	6862	4595	1,493362	7984	4681	1,705618	8074	5099	1,583448
7500x7500	8168	5625	1,452089	9130	5469	1,669409	9819	5662	1,734193
8000x8000	9589	6466	1,482988	10205	6467	1,578011	10723	6425	1,668949
8500x8500	11879	6780	1,752065	11439	7081	1,61545	12478	7130	1,75
9000x9000	12625	7590	1,663373	12540	7577	1,655009	13752	7640	1,8
9500x9500	13567	9029	1,502603	13775	8704	1,582606	15072	8420	1,79
10000x10000	15349	9285	1,653096	15814	9412	1,680195	16675	9213	1,81

Figure 4: Table showing data from 1.40 GHz Quad-Core Intel Core i5 machine.